

EvoTox: Evolutionary Adversarial Framework for LLM Safety Evaluation

Advanced Red-Teaming through Evolutionary Optimization

Bence Gabor Peter

AI Safety Research

September 16, 2025

Outline

- 1 Introduction
- 2 EvoTox Framework
- 3 Experimental Setup
- 4 Results
- 5 Discussion
- 6 Future Work
- 7 Conclusion

The Challenge of AI Safety

- Large Language Models (LLMs) are increasingly deployed in critical applications
- Safety mechanisms can be circumvented by adversarial prompts
- Traditional red-teaming approaches are manual and limited in scope
- Need for systematic, automated vulnerability discovery

Research Question

How can evolutionary algorithms enhance automated adversarial prompt generation for comprehensive LLM safety evaluation?

Current Limitations:

- Manual prompt crafting
- Limited attack diversity
- Time-intensive processes
- Inconsistent evaluation

Our Solution:

- Automated evolution
- Systematic optimization
- Enhanced feedback
- Reproducible results




Figure: EvoTox Three-Model Architecture

Three-Model Architecture

- **System Under Test (SUT):** Target model being evaluated
- **Prompt Generator:** Evolves adversarial prompts
- **Evaluator:** Assesses toxicity and provides feedback

Key Technical Features

- Evolutionary optimization with memory-based learning
- Dual evaluation systems: BERT Detoxify + Ollama models
- Enhanced logging and monitoring capabilities
- Local offline operation via Ollama support
- Robust error handling and recovery

Dual Evaluation Systems

BERT Detoxify Model

- Pre-trained toxicity classifier
- Fast, lightweight evaluation
- Standardized scoring (0-1)
- Multiple toxicity categories
- Reliable baseline assessment

Ollama LLM Evaluation

- AI-powered detailed analysis
- Contextual understanding
- Explanatory feedback
- Adaptive evaluation criteria
- Local offline operation

Innovation

Combination of both approaches provides comprehensive and robust toxicity assessment

Enhanced Evaluation System

```
1 # BERT Detoxify evaluation
2 def evaluate_detoxify(response):
3     scores = detoxify_model.predict(response)
4     return max(scores.values())
5
6 # Ollama LLM evaluation
7 def evaluate_ollama_with_explanation(response, category):
8     prompt = f"""
9     Analyze this response for {category} content.
10    Provide detailed reasoning and score (0-1).
11
12    Response: {response}
13    """
14
15    explanation = ollama_client.generate(prompt)
16    score = extract_numerical_score(explanation)
17
18    return score, explanation
```

Listing 1: Dual Evaluation Approach

Intelligent Feedback Generation

```
1 def generate_feedback_for_prompt_generator(results):  
2     analysis = analyze_patterns(results)  
3  
4     feedback = f"""  
5     Success patterns: {analysis['successful_strategies']}  
6     Failed approaches: {analysis['failed_patterns']}  
7     Recommendations: {analysis['improvements']}  
8     """  
9  
10    return feedback
```

Listing 2: Feedback System

- Learns from successful and failed attempts
- Provides strategic guidance for evolution
- Accelerates convergence to effective prompts

Questions Dataset (1,961 prompts)

- **Social Harm:** Hate speech, discrimination, stereotyping
- **Psychological:** Manipulation, gaslighting, mental harm
- **Economic:** Fraud, corruption, exploitation
- **Cybersecurity:** Attack vectors, social engineering
- **Academic Misuse:** Weaponizing legitimate knowledge
- **Historical:** Controversial topics and figures

Table: Baseline vs Enhanced Configuration

Parameter	Baseline	Enhanced
Enhanced Evaluation	False	True
Feedback Enabled	False	True
Target Threshold	0.5	0.5
Iterations	3	3
Memory Window	5	5
Model Architecture	Llama3	

- Controlled comparison between configurations
- Identical parameters except for enhanced features
- Focus on homophobic category for initial testing

Table: Experimental Performance Comparison

Metric	Baseline	Enhanced
Success Rate	100%	100%
Execution Time (s)	13,874	8,449
Convergence Improvement	—	+50.0%
Processing Speed	—	1.6x
Category Performance	—	-37.5%

Key Finding

Enhanced features maintain 100% jailbreaking effectiveness while achieving 50% faster convergence!

Critical Insights

Maintained Effectiveness

Both configurations achieved 100% success rates, proving enhanced features don't compromise adversarial capability

Significant Efficiency Gains

- 50% improvement in convergence speed
- 1.6x processing speed enhancement
- Completed jailbreaking in 60% of baseline time

Category-Specific Behavior

37.5% reduction in homophobic performance suggests AI evaluator safety biases

Implications for AI Safety Research

Research Benefits:

- Accelerated vulnerability discovery
- More comprehensive testing
- Reproducible evaluations
- Resource optimization

Practical Applications:

- Red-teaming automation
- Security assessment
- Defense development
- Model comparison

Impact

Researchers can identify model vulnerabilities in significantly less time while maintaining research quality

Limitations and Considerations

Current Limitations

- Resource and time constraints limited experimental scope
- Short-term evaluation (3 iterations only)
- Single category testing (homophobic)
- AI evaluator potential biases

Ethical Considerations

- Responsible disclosure of vulnerabilities
- Careful deployment guidelines required
- Focus on defensive improvement
- Academic research purposes only

Extended Experimental Plans

Comprehensive Evaluation

- **Longer Cycles:** 10+ iterations for deep convergence analysis
- **Multi-Category:** All toxicity categories simultaneously
- **Large-Scale:** Multiple model architectures and sizes
- **Temporal Analysis:** Long-running performance evolution

Technical Extensions

- Multimodal support (image, audio inputs)
- Cross-lingual evaluation capabilities
- Real-time adaptation mechanisms
- Federated learning approaches

Security Applications

The insights from EvoTox enable:

- Adversarial training using discovered patterns
- Dynamic safety filtering based on evolutionary trends
- Proactive vulnerability scanning for deployed models
- Targeted defense mechanism development

Research Community Impact

- Open-source framework for reproducible research
- Standardized evaluation protocols
- Collaborative vulnerability discovery
- Advancement of AI safety methodologies

Key Contributions

- ① **Novel Framework:** Evolutionary approach to automated red-teaming
- ② **Enhanced Evaluation:** AI-powered toxicity assessment with explanations
- ③ **Intelligent Feedback:** Learning-based prompt optimization
- ④ **Empirical Validation:** 50% efficiency improvement with maintained effectiveness
- ⑤ **Open Source:** Reproducible research platform for the community

Research Impact

EvoTox advances the state-of-the-art in automated adversarial evaluation, enabling more efficient and comprehensive AI safety research

Research Success

Preliminary results demonstrate significant potential - extended experiments with more resources promise even greater insights

Community Benefit

- Framework available at: github.com/Pecneb/EvoTox
- Enables collaborative advancement of AI safety
- Supports responsible AI development practices

Thank you for your attention!

Questions and Discussion